


## Unix I/O Overview

- Mapping of devices to files allows kernel to export simple interface called *Unix I/O*:
  - Opening and closing files
    - `open()` and `close()`
  - Reading and writing a file
    - `read()` and `write()`
  - Changing the **current file position** (seek)
    - indicates next offset into file to read or write
    - `lseek()`



## Opening Files

- When you open a file →  
Informs the kernel that you are ready to access that file

```
int fd; /* file descriptor */  
  
if ((fd = open("/etc/hosts", O_RDONLY)) < 0) {  
    perror("open");  
    exit(1);  
}
```

- Returns an identifying integer **file descriptor**
  - `fd == -1` indicates that an error occurred
- Each process created by a **Linux shell** begins life with three open files associated with a terminal:
  - 0: standard input (stdin)
  - 1: standard output (stdout)
  - 2: standard error (stderr)

## Reading Files

- Reading a file → copies bytes from the current file position to memory, and then updates file position

```
char buf[512];  
int fd; /* file descriptor */  
int nbytes; /* number of bytes read */  
  
/* Open file fd ... */  
/* Then read up to 512 bytes from file fd */  
if ((nbytes = read(fd, buf, sizeof(buf))) < 0) {  
    perror("read");  
    exit(1);  
}
```

- Returns number of bytes read from file `fd` into `buf`
  - Return type `ssize_t` is signed integer
  - `nbytes < 0` indicates that an error occurred

## File Types

- Each file has a *type* indicating its role in the system
  - *Regular file*: Contains arbitrary data
  - *Directory*: Index for a related group of files
  - *Socket*: For communicating with a process on another machine
- We ignore the other file types (beyond our scope)
  - *Named pipes (FIFOs)*
  - *Symbolic links*
  - *Character and block devices*

## Closing Files

- When you closing a file →  
informs the kernel that you have finished accessing that file

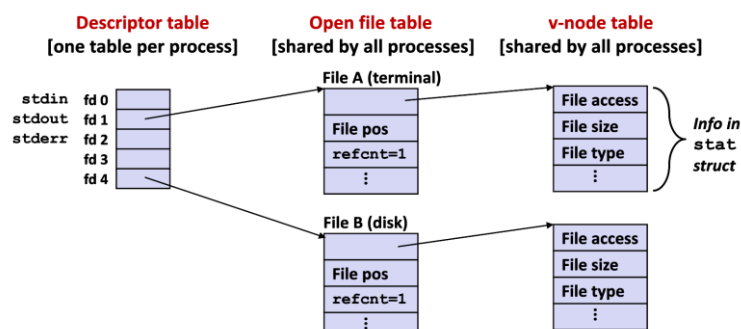
```
int fd; /* file descriptor */  
int retval; /* return value */  
  
if ((retval = close(fd)) < 0) {  
    perror("close");  
    exit(1);  
}
```

## Writing Files

- Writing a file → copies bytes from memory to the current file position, and then updates current file position


```
char buf[512];  
int fd; /* file descriptor */  
int nbytes; /* number of bytes read */  
  
/* Open the file fd ... */  
/* Then write up to 512 bytes from buf to file fd */  
if ((nbytes = write(fd, buf, sizeof(buf))) < 0) {  
    perror("write");  
    exit(1);  
}
```

- Returns number of bytes written from `buf` to file `fd`
  - `nbytes < 0` indicates that an error occurred
- Two descriptors referencing two distinct open files
  - Descriptor 1 (stdout) points to terminal
  - Descriptor 4 points to open disk file



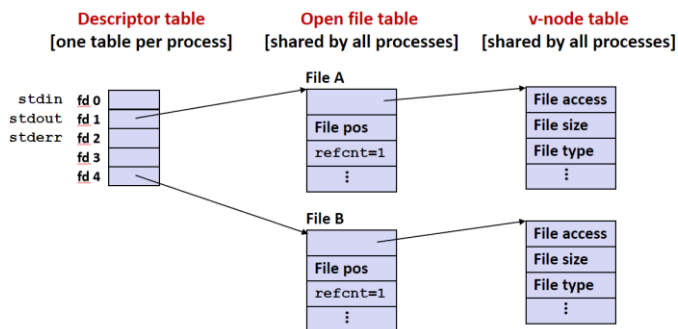
## File Sharing

- Two distinct descriptors sharing the same disk file through two distinct open file table entries
  - E.g., Calling `open` twice with the same `filename` argument



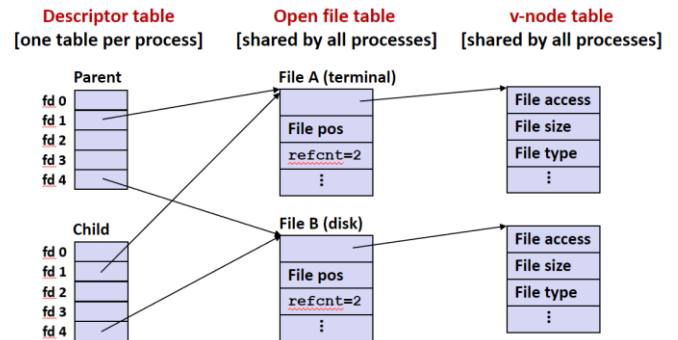
## I/O Redirection Example

- Step #1: open file to which `stdout` should be redirected
  - Happens in child executing shell code, before `exec`



## How Processes Share Files: `fork`

- A child process inherits its parent's open files
- After** `fork`:
  - Child's table same as parent's, and +1 to each `refcnt`



## I/O Redirection Example (cont.)

- Step #2: call `dup2(4, 1)`
  - cause `fd=1` (`stdout`) to refer to disk file pointed at by `fd=4`

